

432 Class 07 Slides

thomaseLove.github.io/432

2021-02-23

Setup

```
library(janitor); library(magrittr)
library(here); library(knitr)

library(skimr)
library(broom)
library(rms)
library(patchwork)
library(GGally)

library(tidyverse)

theme_set(theme_bw())
```

Today's Materials

- The malepstd data
- Using `ols` to fit a linear model
 - Obtaining coefficients and basic summaries
 - Validating summary statistics like R^2
 - ANOVA in `ols`
 - Plot Effects with `summary` and `Predict`
 - Building and using a nomogram
 - Evaluating Calibration
 - Influential points and `dfbeta`
- Spending Degrees of Freedom on Non-Linearity
 - The Spearman ρ^2 (rho-squared) plot
- Building Non-Linear Predictors in `ols`
 - Polynomial Functions
 - Restricted Cubic Splines

The maleptsd data: Background and Exploration

The maleptsd data

The maleptsd file on our web site contains information on PTSD (post traumatic stress disorder) symptoms following childbirth for 64 fathers¹. There are ten predictors and the response is a measure of PTSD symptoms. The raw, untransformed values (ptsd_raw) are right skewed and contain zeros, so we will work with a transformation, specifically, $\text{ptsd} = \log(\text{ptsd_raw} + 1)$ as our outcome, which also contains a lot of zeros.

```
maleptsd <- read_csv(here("data/maleptsd.csv")) %>%  
  clean_names() %>%  
  mutate(ptsd = log(ptsd_raw + 1))
```

¹Source: Ayers et al. 2007 *J Reproductive and Infant Psychology*. The data are described in more detail in Wright DB and London K (2009) *Modern Regression Techniques Using R* Sage Publications.

Skimming the maleptsd data

```
maleptsd %>% select(-id, -ptsd_raw) %>% skim()
```

```
-- Data Summary -----
```

Name	Values
Number of rows	64
Number of columns	11

Column type frequency:	
numeric	11



Group variables	
None	

```
-- Variable type: numeric -----
```

```
# A tibble: 11 x 11
```

	skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
*	<chr>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
1	over2	0	1	2.80	3.34	0	0	1	5	10	
2	over3	0	1	2.72	3.13	0	0	1	5	10	
3	over5	0	1	9.12	1.34	4	9	9.5	10	10	
4	bond	0	1	22.5	3.07	9	21	23	24.2	28	
5	posit	0	1	35.4	11.0	2.5	27.1	37	43.4	50.1	
6	neg	0	1	21.1	11.6	0.7	11.2	20.6	30.4	45.4	
7	contr	0	1	44.2	14.8	4.5	35.1	41.8	54.0	78.5	
8	sup	0	1	13.0	5.87	1.2	9.28	14.2	18.3	20	
9	cons	0	1	49.0	11.2	0	45.8	51	55	65	
10	aff	0	1	8.84	3.08	0	7	9.5	11	17	
11	ptsd	0	1	1.59	1.27	0	0	1.61	2.74	3.95	

Transformation of Outcome



Hmisc::describe() for this transformed outcome

```
maleptsd %$% describe(ptsd)
```

```
ptsd
  n missing distinct      Info      Mean      Gmd      .05      .10
 64      0       24    0.976    1.593    1.458    0.000    0.000
.25    .50    .75    .90    .95
0.000    1.609    2.739    3.246    3.419

lowest : 0.0000000 0.6931472 1.0986123 1.3862944 1.6094379
highest: 3.3322045 3.4339872 3.7612001 3.8286414 3.9512437
```

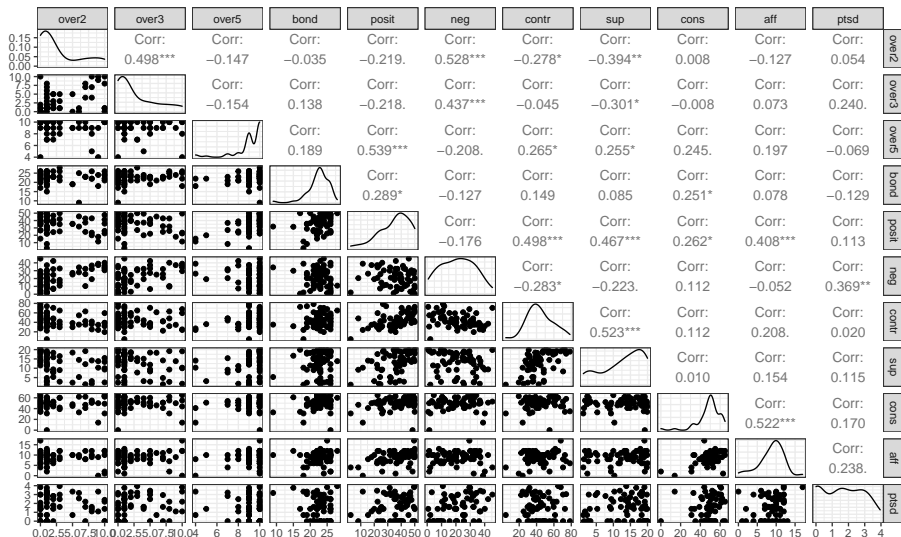
- Gmd is Gini's mean difference, a robust measure of variation.
- If you randomly selected two of the 64 subjects, the average difference in ptsd would be 1.458.

Scatterplot Matrix (code)

```
temp <- maleptsd %>%  
  select(over2, over3, over5, bond, posit, neg, contr,  
         sup, cons, aff, ptsd)  
  
ggpairs(temp)  ## ggpairs from the GGally package
```

- Note that I've placed the outcome (ptsd) last here.
- Result is on the next slide...

Scatterplot Matrix (result)



Using `ols` to fit a linear regression model

Fitting using `ols`

The `ols` function stands for ordinary least squares and comes from the `rms` package, by Frank Harrell and colleagues. Any model fit with `lm` can also be fit with `ols`.

To predict `var_y` using `var_x` from the `my_tibble` data, we would use the following syntax:

```
dd <- datadist(my_tibble)
options(datadist = "dd")

model_name <- ols(var_y ~ var_x, data = my_tibble,
                  x = TRUE, y = TRUE)
```

This leaves the following questions:

- 1 What's the `datadist` stuff doing?
- 2 Why use `x = TRUE`, `y = TRUE` in the fit?

What is datadist?

Before we fit any `ols` model to data from `my_tibble`, we'll use:

```
dd <- datadist(my_tibble)
options(datadist = "dd")
```

Run (the `datadist` code above) once before any models are fitted, storing the distribution summaries for all potential variables. Adjustment values are 0 for binary variables, the most frequent category (or optionally the first category level) for categorical (factor) variables, the middle level for ordered factor variables, and medians for continuous variables.

- excerpted from the `datadist` documentation

Why use `x = TRUE`, `y = TRUE` in the fit?

Once we've set up the distribution summaries with the `datadist` code, we fit linear regression models using the same fitting routines as `lm` with `ols`:

```
model_name <- ols(var_y ~ var_x, data = my_tibble,  
                  x = TRUE, y = TRUE)
```

- `ols` stores additional information beyond what `lm` does
- `x = TRUE` and `y = TRUE` save even more expanded information that we'll need in building plots and summaries of the fit.
- The defaults are `x = FALSE`, `y = FALSE`, but in this class, we'll always want to include these additional pieces.

Using `ols` to fit a Two-Predictor Model

Now, we'll fit an `ols` model predicting our outcome (`ptsd`) using two predictors (`over2` and `over3`) using the `maleptsd` tibble.

- Start with setting the `datadist` up
- Then fit the model, including `x = TRUE`, `y = TRUE`

```
dd <- datadist(maleptsd)
options(datadist = "dd")
```

```
mod_first <- ols(ptsd ~ over2 + over3, data = maleptsd,
                  x = TRUE, y = TRUE)
```

Contents of mod_first?

mod_first

Linear Regression Model

```
ols(formula = ptsd ~ over2 + over3, data = maleptsd, x = TRUE,  
     y = TRUE)
```

		Model Likelihood Ratio Test	Discrimination Indexes
Obs	64	LR chi2 4.18	R2 0.063
sigma	1.2500	d.f. 2	R2 adj 0.033
d.f.	61	Pr(> chi2) 0.1235	g 0.345

Residuals

	Min	1Q	Median	3Q	Max
	-2.259244	-1.337198	0.008866	1.140664	2.333183

	Coef	S.E.	t	Pr(> t)
Intercept	1.3733	0.2202	6.24	<0.0001
over2	-0.0333	0.0544	-0.61	0.5425
over3	0.1149	0.0579	1.98	0.0518

- Likelihood Ratio test?
- What is g?

New elements in `ols`

For our `mod_first`,

- Model Likelihood Ratio test output includes LR $\chi^2 = 4.18$, d.f. = 2, $\text{Pr}(> \chi^2) = 0.1235$

The log of the likelihood ratio, multiplied by -2, yields a test against a χ^2 distribution. Interpret this as a goodness-of-fit test that compares `mod_first` to a null model with only an intercept term. In `ols` this is similar to a global (ANOVA) F test.

- Under the R^2 values, we have $g = 0.345$.
- This is the g -index, based on Gini's mean difference. If you randomly selected two of the subjects in the model, the average difference in predicted `ptsd` will be 0.345.
- This can be compared to the Gini's mean difference for the original `ptsd` values, from `Hmisc::describe`, which was $Gmd = 1.458$.

Validate the summary statistics of an `ols` fit

- Can we validate summary statistics by resampling?

```
set.seed(432010)
validate(mod_first)
```

	index.orig	training	test	optimism	index.corrected	n
R-square	0.0633	0.0812	0.0309	0.0503	0.0130	40
MSE	1.4893	1.4426	1.5408	-0.0982	1.5874	40
g	0.3450	0.3607	0.3055	0.0552	0.2899	40
Intercept	0.0000	0.0000	0.2893	-0.2893	0.2893	40
Slope	1.0000	1.0000	0.8371	0.1629	0.8371	40

- The data used to fit the model provide an over-optimistic view of the quality of fit.
- We're interested here in assessing how well the model might work in new data, and to do so, we can use a resampling approach.
- Consider R^2 here. . .

Interpreting the Resampling Validation Results

	index.orig	training	test	optimism	index.corrected	n
R-sq	0.0633	0.0812	0.0309	0.0503	0.0130	40

- `index.orig` for R^2 is 0.0633. That's what we get from the data we used to fit the model, and is what we see in our standard output.
- With `validate` we create 40 (by default) bootstrapped resamples of the data and then split each of those into training and test samples.
 - For each of the 40 splits, R refits the model (same predictors) in the training sample to obtain R^2 : mean across 40 splits is 0.0812.
 - Check each model in its test sample: average R^2 was 0.0309.
- `optimism` = training result - test result = 0.0503
- `index.corrected` = `index.orig` - `optimism` = 0.0130

While our *nominal* R^2 is 0.0633 for this model, but correcting for optimism yields a *validated* R^2 of 0.0130.

- $R^2 = 0.0130$ better estimates how the model will perform in new data.

ANOVA for mod_first fit by ols

```
anova(mod_first)
```

Analysis of Variance					Response: ptsd	
Factor	d.f.	Partial SS	MS	F	P	
over2	1	0.5862441	0.5862441	0.38	0.5425	
over3	1	6.1458656	6.1458656	3.93	0.0518	
REGRESSION	2	6.4382526	3.2191263	2.06	0.1362	
ERROR	61	95.3127887	1.5625047			

- This adds a line for the complete regression model (both terms) which can be helpful, but is otherwise the same as `anova` after `lm`.
- As with `lm`, this is a sequential ANOVA table, so if we had included `over3` in the model first, we'd get a different SS, MS, F and p for `over2` and `over3`, but the same REGRESSION and ERROR results.

summary for mod_first fit by ols

```
summary(mod_first)
```

Effects				Response : ptsd				
Factor	Low	High	Diff.	Effect	S.E.	Lower 0.95	Upper 0.95	
over2	0	5	5	-0.16658	0.27195	-0.7103800	0.37722	
over3	0	5	5	0.57455	0.28970	-0.0047389	1.15380	

- For over2 effect -0.16658 is the estimated change in ptsd associated with a move from over2 = 0 (see Low value) to over2 = 5 (the High value) assuming no change in over3.
- ols chooses the Low and High values from the interquartile range.


```
maleptsd %$% quantile(over2, c(0.25, 0.75))
```

```
25% 75%  
0    5
```

Plot the summary to see effect sizes

- Goal: plot effect sizes for similar moves within predictor distributions.


```
plot(summary(mod_first))
```



- The triangles indicate the point estimate, augmented with confidence interval bars.
 - The 90% confidence intervals are plotted with the thickest bars.
 - The 95% CIs are then shown with thinner, more transparent bars.
 - Finally, the 99% CIs are shown as the longest, thinnest bars.

What do the individual effects look like?


```
ggplot(Predict(mod_first, conf.int = 0.95), layout = c(1,2))
```



- The left plot shows the impact of changing over2 on ptsd holding over3 constant at the median over3 value, which is 1.
- The right plot shows the impact of changing over3 on ptsd holding over2 constant at its median value, which is 1.
- Defaults: add 95% CI bands and layout tries for a square.

Build a nomogram for the `ols` fit

```
plot(nomogram(mod_first))
```



Nomograms


For complex models (this model isn't actually very complex) it can be helpful to have a tool that will help you see the effects of model in terms of their impact on the predicted outcome.

A *nomogram* is an established graphical tool for doing this.

- Find the value of each predictor on its provided line, and identify the “points” for that predictor by drawing a vertical line up to the “Points”.
- Then sum up the points over all predictors to obtain “Total Points”.
- Draw a vertical line down from the “Total Points” to the “Linear Predictor” to get the predicted ptsd for this subject.

Using the nomogram for the o1s fit

Predicted ptsd for a subject with $\text{over2} = 8$ and $\text{over3} = 7$?



Actual Prediction for such a subject...

- The `predict` function for an `ols` fit can provide fitted values.

```
predict(mod_first,  
        newdata = tibble(over2 = 8, over3 = 7))
```

1

1.911123

- The `broom` package doesn't (really) support `rms` fits, and throws a warning (omitted here), but you could always refit the model with `lm...`

```
augment(mod_first,  
        newdata = tibble(over2 = 8, over3 = 7))
```

```
# A tibble: 1 x 3  
  over2 over3 .fitted  
  <dbl> <dbl>   <dbl>  
1     8     7     1.91
```

Assessing the Calibration of `mod_first`

We would like our model to be well-calibrated, in the following sense...


- Suppose our model assigns a predicted outcome of 6 to several subjects. If the model is well-calibrated, then we expect the mean of those subjects' actual outcomes to be very close to 6.

We'd like to look at the relationship between the observed `ptsd` outcome and our predicted `ptsd` from the model.

- The calibration plot we'll create provides two estimates (with and without bias-correction) of the predicted vs. observed values of our outcome, and compares these to the ideal scenario (predicted = observed).
- The plot uses resampling validation to produce bias-corrected estimates and uses lowess smooths to connect across predicted values.
- Calibration plots require `x = TRUE`, `y = TRUE` in the `ols` fit.

Calibration Plot for mod_first

```
set.seed(432); plot(calibrate(mod_first))
```



n=64 Mean absolute error=0.224 Mean squared error=0.1048
0.9 Quantile of absolute error=0.64

Influential Points for `mod_first`?

The `dfbeta` value for a particular subject and coefficient β is the change in the coefficient that happens when the subject is excluded from the model.

```
which.influence(mod_first, cutoff = 0.3)
```

```
$over2
```

```
[1] 32 33 37 40 43 57
```

```
$over3
```

```
[1] 32 43 57
```

- These are the subjects that have absolute values of `dfbetas` that exceed the specified cutoff (default is 0.2)

Show the influential points more directly?


```
w <- which.influence(mod_first, cutoff = 0.3)
d <- maleptsd %>% select(over2,over3,ptsd) %>% data.frame()
show.influence(w, d)
```

	Count	over2	over3
32	2	* 0	*10
33	1	*10	1
37	1	*10	0
40	1	*10	1
43	2	* 6	* 0
57	2	* 1	* 8

- Count = number of coefficients where this row appears influential.
- Use `maleptsd %>% slice(32)` to see row 32 in its entirety.
- Use residual plots (with an `lm` fit) to check Cook's distances.

Residuals vs. Fitted Values is an easy plot

```
plot(resid(mod_first) ~ fitted(mod_first))
```



Fitting all Residual Plots for mod_first





To fit more complete residual plots (and to do other things) we will fit the `lm` version of this same model...

```
mod_first_lm <- lm(ptsd ~ over2 + over3, data = maleptsd)

par(mfrow = c(2,2))
plot(mod_first_lm)
par(mfrow = c(1,1))
```

- Plots are shown on the next slide. While the subject in row 32 is more influential than most other points, it doesn't reach the standard of a problematic Cook's distance.

Residual Plots for mod_first



Thinking about Non-Linear Terms?

Non-Linear Terms

In building a linear regression model, we're most often going to be thinking about:

- for quantitative predictors, some curvature. . .
 - perhaps polynomial terms
 - but more often restricted cubic splines
- for any predictors, possible interactions
 - between categorical predictors
 - between categorical and quantitative predictors
 - between quantitative predictors

Polynomial Regression

A polynomial in the variable x of degree D is a linear combination of the powers of x up to D . Fitting such a model creates a **polynomial regression**.

- Linear: $y = \beta_0 + \beta_1 x$
- Quadratic: $y = \beta_0 + \beta_1 x + \beta_2 x^2$
- Cubic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3$
- Quartic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4$
- Quintic: $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \beta_5 x^5$

An **orthogonal polynomial** sets up a model design matrix and then scales those columns so that each column is uncorrelated with the previous ones.

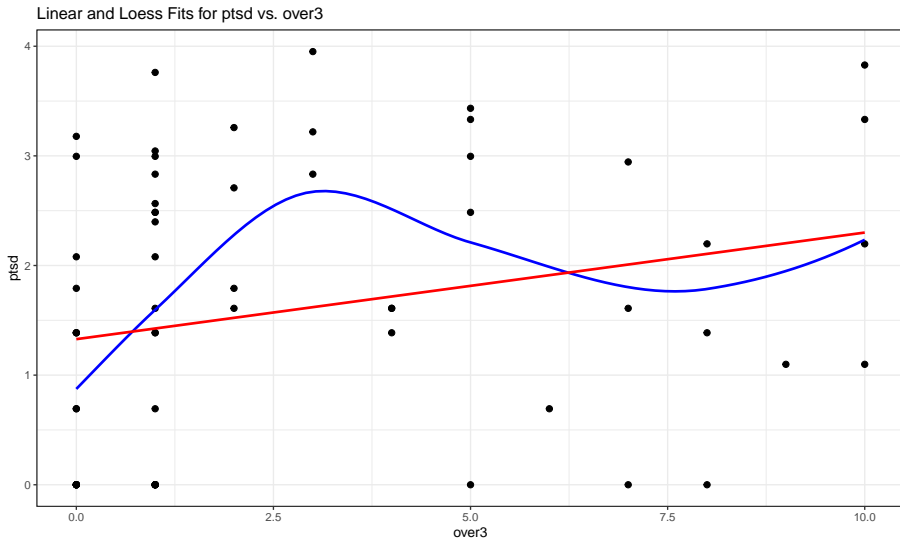
- This reduction in collinearity (correlation between predictors) lets us gauge whether the addition of any particular polynomial term improves model fit.

Using over3 to predict ptsd

- Let's look at both a linear fit and a loess smooth to see if they indicate meaningfully different things about the association.

```
ggplot(maleptsd, aes(x = over3, y = ptsd)) +  
  geom_point(size = 2) +  
  geom_smooth(method = "loess", formula = y ~ x,  
              se = FALSE, col = "blue") +  
  geom_smooth(method = "lm", formula = y ~ x,  
              se = FALSE, col = "red") +  
  labs(title = "Linear and Loess Fits for ptsd vs. over3")
```

Linear and Loess Fits for ptsd with over3



Fitting polynomial regressions with `ols`

```
dd <- datadist(maleptsd)
options(datadist = "dd")

mod_B1 <- ols(ptsd ~ over3,
              data = maleptsd, x = TRUE, y = TRUE)
mod_B2 <- ols(ptsd ~ pol(over3, 2),
              data = maleptsd, x = TRUE, y = TRUE)
mod_B3 <- ols(ptsd ~ pol(over3, 3),
              data = maleptsd, x = TRUE, y = TRUE)
```

- Note the use of `pol()` from the `rms` package here to fit orthogonal polynomials, rather than `poly()` which we used for an `lm` fit.

Model B1 (linear in over3)

mod_B1

Linear Regression Model

```
ols(formula = ptsd ~ over3, data = maleptsd, x = TRUE, y = TRUE)
```

		Model Likelihood		Discrimination	
		Ratio Test		Indexes	
Obs	64	LR chi2	3.79	R2	0.058
sigma1.2437		d.f.	1	R2 adj	0.042
d.f.	62	Pr(> chi2)	0.0515	g	0.319

Residuals

Min	1Q	Median	3Q	Max
-2.106040	-1.328148	0.009528	1.079533	2.335816

	Coef	S.E.	t	Pr(> t)
Intercept	1.3281	0.2065	6.43	<0.0001
over3	0.0972	0.0500	1.95	0.0563

Model B2 (quadratic polynomial in over3)

mod_B2

Linear Regression Model

```
ols(formula = ptsd ~ pol(over3, 2), data = maleptsd, x = TRUE,  
     y = TRUE)
```

		Model Likelihood Ratio Test	Discrimination Indexes
Obs	64	LR chi2 6.26	R2 0.093
sigma	1.2299	d.f. 2	R2 adj 0.063
d.f.	61	Pr(> chi2) 0.0437	g 0.437

Residuals

	Min	1Q	Median	3Q	Max
	-2.186161	-1.121078	0.005005	1.000865	2.311934

	Coef	S.E.	t	Pr(> t)
Intercept	1.1211	0.2440	4.59	<0.0001
over3	0.3575	0.1751	2.04	0.0455
over3^2	-0.0293	0.0189	-1.55	0.1264

Model B3 (cubic polynomial in over3)

mod_B3

Linear Regression Model

```
ols(formula = ptsd ~ pol(over3, 3), data = maleptsd, x = TRUE,  
     y = TRUE)
```

		Model Likelihood	Discrimination
		Ratio Test	Indexes
Obs	64	LR chi2 13.92	R2 0.195
sigma	1.1681	d.f. 3	R2 adj 0.155
d.f.	60	Pr(> chi2) 0.0030	g 0.622

Residuals

	Min	1Q	Median	3Q	Max
	-1.96266	-0.77151	-0.07836	0.88211	2.40654

	Coef	S.E.	t	Pr(> t)
Intercept	0.7715	0.2641	2.92	0.0049
over3	1.2292	0.3568	3.44	0.0010
over3^2	-0.2900	0.0961	-3.02	0.0037
over3^3	0.0184	0.0067	2.76	0.0076

Store the polynomial fits

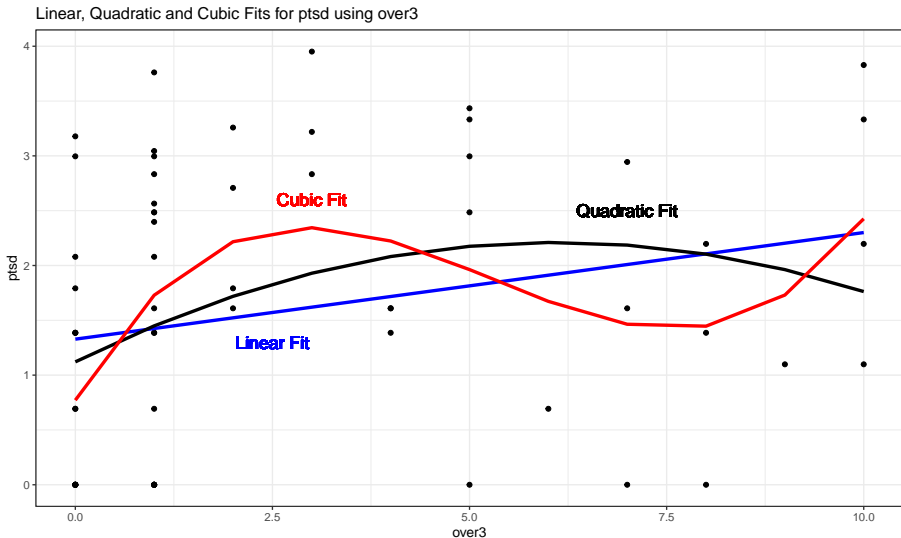
First, we need to store the values. Again broom doesn't play well with `ols` fits, so I'll just add the predictions as columns

```
ptsd_fits <- maleptsd %>%  
  mutate(fitB1 = predict(mod_B1),  
         fitB2 = predict(mod_B2),  
         fitB3 = predict(mod_B3))
```

Code to plot polynomial fits

```
ggplot(ptsd_fits, aes(x = over3, y = ptsd)) +  
  geom_point() +  
  geom_line(aes(x = over3, y = fitB1),  
            col = "blue", size = 1.25) +  
  geom_line(aes(x = over3, y = fitB2),  
            col = "black", size = 1.25) +  
  geom_line(aes(x = over3, y = fitB3),  
            col = "red", size = 1.25) +  
  geom_text(x = 2.5, y = 1.3, label = "Linear Fit",  
            size = 5, col = "blue") +  
  geom_text(x = 7, y = 2.5, label = "Quadratic Fit",  
            size = 5, col = "black") +  
  geom_text(x = 3, y = 2.6, label = "Cubic Fit",  
            size = 5, col = "red") +  
  labs(title = "Linear, Quadratic and Cubic Fits for ptsd us
```

The Polynomial Fits, plotted

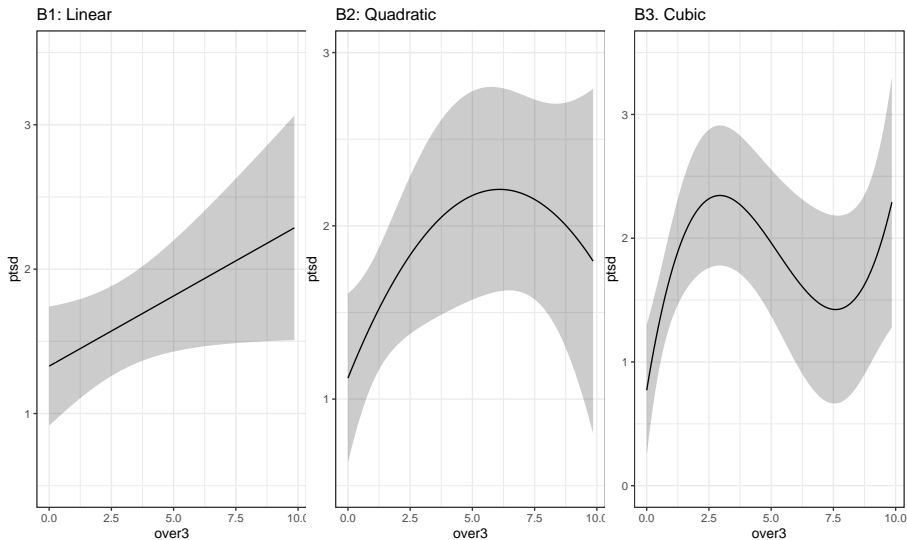


Code to plot polynomial fits with Predict

```
p1 <- ggplot(Predict(mod_B1)) + ggtitle("B1: Linear")
p2 <- ggplot(Predict(mod_B2)) + ggtitle("B2: Quadratic")
p3 <- ggplot(Predict(mod_B3)) + ggtitle("B3. Cubic")

p1 + p2 + p3
```

Visualizing the polynomial fits with Predict



Splines

- A **linear spline** is a continuous function formed by connecting points (called **knots** of the spline) by line segments.
- A **restricted cubic spline** is a way to build highly complicated curves into a regression equation in a fairly easily structured way.
- A restricted cubic spline is a series of polynomial functions joined together at the knots.
 - Such a spline gives us a way to flexibly account for non-linearity without over-fitting the model.
 - Restricted cubic splines can fit many different types of non-linearities.
 - Specifying the number of knots is all you need to do in R to get a reasonable result from a restricted cubic spline.

The most common choices are 3, 4, or 5 knots.

- 3 Knots, 2 degrees of freedom, allows the curve to “bend” once.
- 4 Knots, 3 degrees of freedom, lets the curve “bend” twice.
- 5 Knots, 4 degrees of freedom, lets the curve “bend” three times.

Fitting Restricted Cubic Splines with `ols`

Let's consider a restricted cubic spline model for `ptsd` based on `over3` with:

- 3 knots in `modC3`, 4 knots in `modC4`, and 5 knots in `modC5`

```
dd <- datadist(maleptsd)
options(datadist = "dd")

mod_C3 <- ols(ptsd ~ rcs(over3, 3),
              data = maleptsd, x = TRUE, y = TRUE)
mod_C4 <- ols(ptsd ~ rcs(over3, 4),
              data = maleptsd, x = TRUE, y = TRUE)
mod_C5 <- ols(ptsd ~ rcs(over3, 5),
              data = maleptsd, x = TRUE, y = TRUE)
```

Model C3 (3-knot spline in over3)

mod_C3

Linear Regression Model

```
ols(formula = ptsd ~ rcs(over3, 3), data = maleptsd, x = TRUE,  
     y = TRUE)
```

		Model Likelihood Ratio Test	Discrimination Indexes
Obs	64	LR chi2 8.14	R2 0.119
sigma	1.2119	d.f. 2	R2 adj 0.091
d.f.	61	Pr(> chi2) 0.0171	g 0.497

Residuals

	Min	1Q	Median	3Q	Max
	-2.25280	-1.01211	-0.08417	0.99283	2.26912

	Coef	S.E.	t	Pr(> t)
Intercept	1.0121	0.2525	4.01	0.0002
over3	0.5041	0.2024	2.49	0.0155
over3'	-1.9561	0.9444	-2.07	0.0426

Model C4 (4-knot spline in over3)

mod_C4

Linear Regression Model

```
ols(formula = ptsd ~ rcs(over3, 4), data = maleptsd, x = TRUE,  
     y = TRUE)
```

		Model Likelihood Ratio Test	Discrimination Indexes
Obs	64	LR chi2 9.60	R2 0.139
sigma	1.2081	d.f. 3	R2 adj 0.096
d.f.	60	Pr(> chi2) 0.0223	g 0.498

Residuals

	Min	1Q	Median	3Q	Max
	-2.0195	-0.8429	-0.1498	0.9543	2.3351

	Coef	S.E.	t	Pr(> t)
Intercept	0.8429	0.2898	2.91	0.0051
over3	0.9535	0.4316	2.21	0.0310
over3'	-8.9140	5.9823	-1.49	0.1414
over3''	13.4801	9.6354	1.40	0.1670

Model C5 (5-knot spline in over3)

mod_C5

Linear Regression Model

```
ols(formula = ptsd ~ rcs(over3, 5), data = maleptsd, x = TRUE,  
     y = TRUE)
```

		Model Likelihood Ratio Test	Discrimination Indexes
Obs	64	LR chi2 10.58	R2 0.152
sigma	1.1989	d.f. 3	R2 adj 0.110
d.f.	60	Pr(> chi2) 0.0142	g 0.520

Residuals

Min	1Q	Median	3Q	Max
-1.9251	-0.8060	-0.1128	1.0032	2.3721

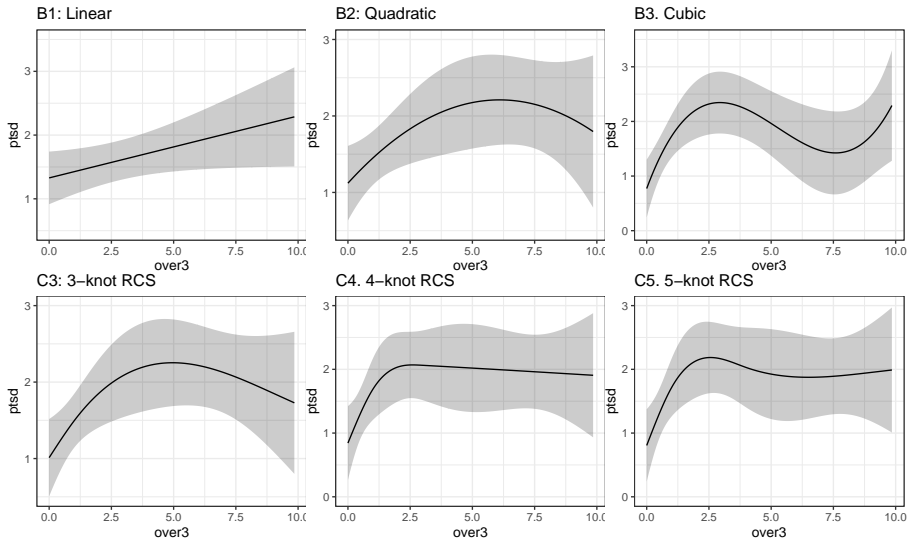
	Coef	S.E.	t	Pr(> t)
Intercept	0.8060	0.2839	2.84	0.0062
over3	1.0137	0.3889	2.61	0.0115
over3'	-8.5106	4.3895	-1.94	0.0572
over3''	11.6766	6.2891	1.86	0.0683

Code to plot all six fits

```
p1 <- ggplot(Predict(mod_B1)) + ggtitle("B1: Linear")
p2 <- ggplot(Predict(mod_B2)) + ggtitle("B2: Quadratic")
p3 <- ggplot(Predict(mod_B3)) + ggtitle("B3: Cubic")
p4 <- ggplot(Predict(mod_C3)) + ggtitle("C3: 3-knot RCS")
p5 <- ggplot(Predict(mod_C4)) + ggtitle("C4: 4-knot RCS")
p6 <- ggplot(Predict(mod_C5)) + ggtitle("C5: 5-knot RCS")

(p1 + p2 + p3) / (p4 + p5 + p6)
```

Visualizing the fits better?



Data Spending: Non-Linearity Prior to Fits

Spending degrees of freedom wisely

- Suppose we have a data set with many possible predictors, and minimal theory or subject matter knowledge to guide us.
- We might want our final inferences to be as unbiased as possible. To accomplish this, we have to pay a penalty (in terms of degrees of freedom) for any “peeks” we make at the data in advance of fitting a model.
- So that rules out a lot of decision-making about non-linearity based on looking at the data, if our sample size isn't much larger than 15 times the number of predictors we're considering including in our model.
- In our case, we have $n = 64$ observations on 10 predictors.
- In addition, adding non-linearity to our model costs additional degrees of freedom.
- What can we do?

Spearman's ρ^2 plot: A smart first step?


Spearman's ρ^2 is an indicator (not a perfect one) of potential predictive punch, but doesn't give away the game.

- Idea: Perhaps we should focus our efforts re: non-linearity on predictors that score better on this measure.

```
spear_ptsd <- spearman2(ptsd ~ over2 + over3 + over5 + bond +  
                        posit + neg + contr + sup + cons + aff,  
                        data = maleptsd)
```

Spearman's ρ^2 Plot

```
plot(spear_ptsd)
```



Conclusions from Spearman ρ^2 Plot

- `neg` is the most attractive candidate for a non-linear term, as it packs the most potential predictive punch, so if it does turn out to need non-linear terms, our degrees of freedom will be well spent.
 - By no means is this suggesting that `neg` actually needs a non-linear term, or will show significant non-linearity. We'd have to fit a model with and without non-linearity in `neg` to know that.
 - Non-linearity will often take the form of a product term, a polynomial term, or a restricted cubic spline.
 - Since all of these predictors are quantitative, we'll think about polynomial or spline terms, soon.
- `over3`, also quantitative, has the next most potential predictive punch
- these are followed by `cons` and `aff`

With 64 observations (63 df) we should be thinking about models with relatively tiny numbers of regression inputs.

- Non-linear terms (polynomials, splines) just add to the problem, as they need additional df to be estimated.

In this case, we might choose between

- including non-linearity in one (or maybe 2) variables (and that's it),
- or a linear model including maybe 3-4 predictors, tops

and even that would be tough to justify with this small sample size.

Contents of spear_ptsd

```
spear_ptsd
```

```
Spearman rho^2      Response variable:ptsd
```

	rho2	F	df1	df2	P	Adjusted rho2	n
over2	0.021	1.34	1	62	0.2522	0.005	64
over3	0.110	7.67	1	62	0.0074	0.096	64
over5	0.006	0.36	1	62	0.5527	-0.010	64
bond	0.004	0.22	1	62	0.6405	-0.013	64
posit	0.008	0.50	1	62	0.4825	-0.008	64
neg	0.136	9.73	1	62	0.0027	0.122	64
contr	0.001	0.03	1	62	0.8602	-0.016	64
sup	0.004	0.23	1	62	0.6357	-0.012	64
cons	0.052	3.40	1	62	0.0699	0.037	64
aff	0.052	3.39	1	62	0.0704	0.037	64

Proposed New Model

Fit a model to predict ptsd using:

- a 4-knot spline on neg
- a 3-knot spline on over3
- a linear term on cons
- a linear term on aff

Perhaps more than we can reasonably do with 64 observations, but let's see how it looks.

```
dd <- datadist(maleptsd)
options(datadist = "dd")
```

```
mod_second <- ols(ptsd ~ rcs(neg, 4) + rcs(over3, 3) +
                  cons + aff, data = maleptsd,
                  x = TRUE, y = TRUE)
```

Our second model

mod_second

Linear Regression Model

```
ols(formula = ptsd ~ rcs(neg, 4) + rcs(over3, 3) + cons + aff,  
     data = maleptsd)
```

		Model Likelihood Ratio Test	Discrimination Indexes
Obs	64	LR chi2 21.28	R2 0.283
sigma	1.1415	d.f. 7	R2 adj 0.193
d.f.	56	Pr(> chi2) 0.0034	g 0.763

Residuals

Min	1Q	Median	3Q	Max
-2.06529	-0.81434	0.06745	0.81760	2.17200

	Coef	S.E.	t	Pr(> t)
Intercept	-0.4255	0.7490	-0.57	0.5723
neg	0.0660	0.0603	1.10	0.2780
neg'	-0.1261	0.1641	-0.77	0.4456
neg''	0.4924	0.5373	0.92	0.3634
over3	0.4582	0.2007	2.28	0.0263
over3'	-2.1247	0.9433	-2.25	0.0282
cons	-0.0119	0.0164	-0.72	0.4722
aff	0.1450	0.0598	2.42	0.0186

ANOVA for this model

```
anova(mod_second)
```

Analysis of Variance

Response: ptsd

Factor	d.f.	Partial SS	MS	F	P
neg	3	11.4062336	3.8020779	2.92	0.0420
Nonlinear	2	1.6536591	0.8268295	0.63	0.5339
over3	2	6.8378486	3.4189243	2.62	0.0814
Nonlinear	1	6.6106843	6.6106843	5.07	0.0282
cons	1	0.6826901	0.6826901	0.52	0.4722
aff	1	7.6565797	7.6565797	5.88	0.0186
TOTAL NONLINEAR	3	7.8079300	2.6026433	2.00	0.1248
REGRESSION	7	28.7821644	4.1117378	3.16	0.0070
ERROR	56	72.9688769	1.3030157		

- Remember that this ANOVA testing is sequential.

Validation of Summary Statistics

```
set.seed(432); validate(mod_second)
```

	index.orig	training	test	optimism	index.corrected	n
R-square	0.2829	0.3725	0.1566	0.2159	0.0670	40
MSE	1.1401	0.9734	1.3409	-0.3675	1.5076	40
g	0.7630	0.8562	0.6503	0.2059	0.5571	40
Intercept	0.0000	0.0000	0.4564	-0.4564	0.4564	40
Slope	1.0000	1.0000	0.7402	0.2598	0.7402	40


summary results for mod_second

```
summary(mod_second)
```

Effects				Response : ptsd				
Factor	Low	High	Diff.	Effect	S.E.	Lower 0.95	Upper 0.95	
neg	11.175	30.425	19.25	0.44727	0.41704	-0.388160	1.28270	
over3	0.000	5.000	5.00	0.90059	0.43913	0.020902	1.78030	
cons	45.750	55.000	9.25	-0.10997	0.15192	-0.414310	0.19437	
aff	7.000	11.000	4.00	0.57998	0.23926	0.100680	1.05930	


Plot of summary results for mod_second

```
plot(summary(mod_second))
```




Nomogram for mod_second

```
plot(nomogram(mod_second))
```




Seeing the impact of the modeling another way

```
ggplot(Predict(mod_second))
```



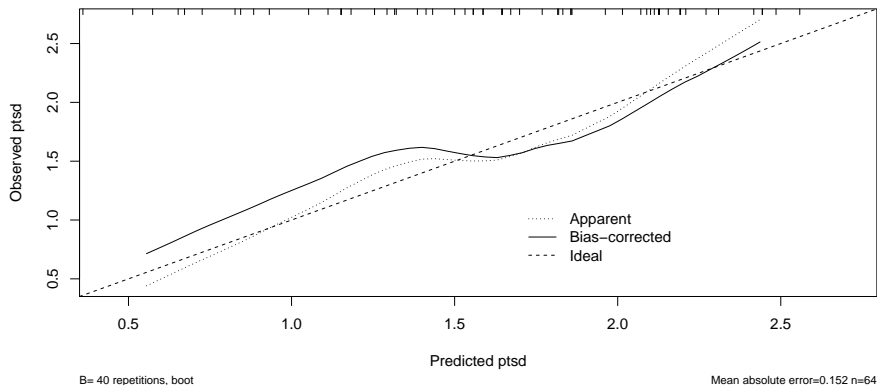
Residuals vs. Fitted Values to check assumptions

```
plot(resid(mod_second) ~ fitted(mod_second))
```



Checking the model's calibration

```
set.seed(432); plot(calibrate(mod_second))
```



n=64 Mean absolute error=0.152 Mean squared error=0.03106
0.9 Quantile of absolute error=0.279

Limitations of `lm` for fitting complex linear models

We can certainly assess this big, complex model using `lm`, too:

- with in-sample summary statistics like adjusted R^2 , AIC and BIC,
- we can assess its assumptions with residual plots, and
- we can also compare out-of-sample predictive quality through cross-validation,

But to really delve into the details of how well this complex model works, and to help plot what is actually being fit, we'll probably want to fit the model using `ols`.

- In Project 1, we expect some results that are most easily obtained using `lm` and others that are most easily obtained using `ols`.

Next Time

- The HERS data
- Fitting a more complex linear regression model
 - Dealing with categorical predictors
 - Dealing with interactions (another form of non-linearity)
 - Adding missing data into all of this, and running multiple imputation

Don't forget to complete the Minute Paper after Class 7 by tomorrow at Noon!